

Change of Variables and Normalizing Flows

Cheng Soon Ong
Marc Peter Deisenroth

December 2020



Normalizing flows for density estimation




Figure: Generated with PyMC3 (Salvatier et al., 2016)

Key idea

(Tabak & Turner, 2013; Rippel & Adams, 2013; Rezende & Mohamed, 2015)

Build complex distributions from simple distributions via a flow of successive (invertible) transformations

Normalizing flows for density estimation




Figure: Generated with PyMC3 (Salvatier et al., 2016)

Key idea


(Tabak & Turner, 2013; Rippel & Adams, 2013; Rezende & Mohamed, 2015)

Build complex distributions from simple distributions via a flow of successive (invertible) transformations

Key ingredient: **Change-of-variables trick**

Change of Variables


Change of variables: Key idea



Key idea

Transform random variable X into random variable Z using an invertible transformation ϕ , while keeping track of the change in distribution

Change of variables: Key idea




Key idea

Transform random variable X into random variable Z using an invertible transformation ϕ , while keeping track of the change in distribution

- ▶ Distribution p_X induces distribution p_Z via transformation ϕ
- ▶ Distribution p_Z induces distribution p_X via transformation ϕ^{-1}

Jacobian determinant



- ▶ Determinant of Jacobian


$$\left| \det \left(\frac{dz}{dx} \right) \right| = \left| \det \left(\frac{d\phi(x)}{dx} \right) \right|$$

tells us how much the domain dx is stretched to dz

How it works

- ▶ Constraint: volume preservation


$$\int_{\mathcal{X}} p_X(\mathbf{x}) d\mathbf{x} = 1 = \int_{\mathcal{Z}} p_Z(\mathbf{z}) d\mathbf{z}$$



How it works


- ▶ Constraint: volume preservation

$$\int_{\mathcal{X}} p_X(\mathbf{x}) d\mathbf{x} = 1 = \int_{\mathcal{Z}} p_Z(\mathbf{z}) d\mathbf{z}$$




- ▶ Volume preservation: Rescale p_Z by the inverse of Jacobian determinant

$$p_Z(\mathbf{z}) = p_X(\mathbf{x}) \left| \det \left(\frac{d\phi(\mathbf{x})}{d\mathbf{x}} \right) \right|^{-1}$$



Considerations



$$p_Z(z) = p_X(x) \left| \det \left(\frac{d\phi(x)}{dx} \right) \right|^{-1}$$

- ▶ Express target distribution p_Z in terms of known distribution p_X and the Jacobian determinant of an invertible mapping ϕ
- ▶ No need to invert ϕ explicitly
- ▶ Generate expressive distributions p_Z by simple p_X and flexible transformation ϕ

Applications

- ▶ Numerical integration (turn indefinite integrals into definite ones)
- ▶ Neural ODEs (E 2017, Chen et al., 2018)
- ▶ Learning in implicit generative models (e.g., GANs) and likelihood-free inference (e.g., ABC)
(e.g., Mohamed & Lakshminarayanan, 2016; Sisson et al., 2007)
- ▶ **Normalizing flows** (Rezende & Mohamed, 2015)

Normalizing Flows

Normalizing flows for density estimation





Figure: Generated with PyMC3 (Salvatier et al., 2016)

Key idea

(Tabak & Turner, 2013; Rippel & Adams, 2013; Rezende & Mohamed, 2015)


Build complex distributions from simple distributions via a flow of successive (invertible) transformations

How it works



- ▶ Random variable $z_0 \sim p_0$
- ▶ Simple base distribution p_0 , e.g. $p_0 = \mathcal{N}(\mathbf{0}, \mathbf{I})$

How it works




- ▶ Random variable $z_0 \sim p_0$
- ▶ Simple base distribution p_0 , e.g. $p_0 = \mathcal{N}(\mathbf{0}, \mathbf{I})$
- ▶ Successive transformation of z_k via invertible transformations f_k :

$$z_k = f_k(z_{k-1})$$

- ▶ Observed data $x = z_K$ at the end of the chain

$$x = z_K = f_K \circ f_{K-1} \circ \cdots \circ f_1(z_0)$$

Marginal distribution



- ▶ Repeated application of **change-of-variables** trick

$$p(\mathbf{x}) = p(\mathbf{z}_K) = p(\mathbf{z}_0) \prod_{k=1}^K \left| \det \frac{df_k(\mathbf{z}_{k-1})}{d\mathbf{z}_{k-1}} \right|^{-1}$$

- ▶ Entropy

$$\log p(\mathbf{x}) = \log p(\mathbf{z}_K) = \log p(\mathbf{z}_0) - \sum_{k=1}^K \log \left| \det \left(\frac{df_k(\mathbf{z}_{k-1})}{d\mathbf{z}_{k-1}} \right) \right|$$





Illustration with PyMC3 (Salvatier et al., 2016)




1 planar flow


Figure: Generated using a PyMC3 tutorial (Salvatier et al., 2016)

- ▶ Repeated application of a planar flow $\mathbf{z}_k = f_k(\mathbf{z}_{k-1}) = \mathbf{z}_{k-1} + \mathbf{u}\sigma(\mathbf{w}^\top \mathbf{z}_{k-1} + b)$

Illustration with PyMC3 (Salvatier et al., 2016)



1 planar flow



2 planar flows

Figure: Generated using a PyMC3 tutorial (Salvatier et al., 2016)

- ▶ Repeated application of a planar flow $\mathbf{z}_k = f_k(\mathbf{z}_{k-1}) = \mathbf{z}_{k-1} + \mathbf{u}\sigma(\mathbf{w}^\top \mathbf{z}_{k-1} + b)$

Illustration with PyMC3 (Salvatier et al., 2016)






Figure: Generated using a PyMC3 tutorial (Salvatier et al., 2016)

- ▶ Repeated application of a planar flow $\mathbf{z}_k = f_k(\mathbf{z}_{k-1}) = \mathbf{z}_{k-1} + \mathbf{u}\sigma(\mathbf{w}^\top \mathbf{z}_{k-1} + b)$


Illustration with PyMC3 (Salvatier et al., 2016)




1 planar flow



2 planar flows



3 planar flows



7 planar flows

Figure: Generated using a PyMC3 tutorial (Salvatier et al., 2016)

- ▶ Repeated application of a planar flow $\mathbf{z}_k = f_k(\mathbf{z}_{k-1}) = \mathbf{z}_{k-1} + \mathbf{u}\sigma(\mathbf{w}^\top \mathbf{z}_{k-1} + b)$

Illustration with PyMC3 (Salvatier et al., 2016)




Figure: Generated using a PyMC3 tutorial (Salvatier et al., 2016)

- ▶ Repeated application of a planar flow $\mathbf{z}_k = f_k(\mathbf{z}_{k-1}) = \mathbf{z}_{k-1} + \mathbf{u}\sigma(\mathbf{w}^\top \mathbf{z}_{k-1} + b)$

Computing expectations

$$\mathbb{E}_{p_X}[l(\mathbf{x})] = \mathbb{E}_{p_K}[l(\mathbf{z}_K)] = \mathbb{E}_{p_0}[l(f_K \circ \dots \circ f_1(\mathbf{z}_0))]$$

Computing expectations

$$\mathbb{E}_{p_X}[l(\mathbf{x})] = \mathbb{E}_{p_K}[l(\mathbf{z}_K)] = \mathbb{E}_{p_0}[l(f_K \circ \dots \circ f_1(\mathbf{z}_0))]$$

- ▶ Expectations w.r.t. p_K can be computed without explicitly knowing p_K or p_X
 - ▶ Sample $\mathbf{z}_0^{(s)} \sim p_0$
 - ▶ Push sample forward through sequence of deterministic transformations
 - ▶▶ Valid sample $\mathbf{x}^{(s)} \sim p_X(x)$
- ▶ Monte-Carlo estimation to get expected value

Computational considerations

- ▶ Compute log-determinant of Jacobian
- ▶ Cheap (linear) if Jacobian is (block-)diagonal or triangular

Computational considerations

- ▶ Compute log-determinant of Jacobian
- ▶ Cheap (linear) if Jacobian is (block-)diagonal or triangular
- ▶ Require partial derivatives

$$\frac{\partial z_k^{(d)}}{\partial z_{k-1}^{(>d)}} = 0 \quad \Rightarrow \quad \frac{dz_k}{dz_{k-1}} = \begin{bmatrix} \frac{\partial z_k^{(1)}}{\partial z_{k-1}^{(1)}} & \mathbf{0} & \dots & \mathbf{0} \\ \vdots & \frac{\partial z_k^{(2)}}{\partial z_{k-1}^{(2)}} & \dots & \vdots \\ \frac{\partial z_k^{(D)}}{\partial z_{k-1}^{(1)}} & \dots & \dots & \frac{\partial z_k^{(D)}}{\partial z_{k-1}^{(D)}} \end{bmatrix} \in \mathbb{R}^{D \times D}.$$

Autoregressive flows

- ▶ High-level idea:

$$z_k^{(d)} = \phi(z_{k-1}^{(\leq d)})$$

Autoregressive flows

- ▶ High-level idea:

$$z_k^{(d)} = \phi(z_{k-1}^{(\leq d)})$$


- ▶ NICE (Dinh et al., 2014)
- ▶ Inverse autoregressive flow (Kingma et al., 2016)
- ▶ Real NVP (Dinh et al., 2017)
- ▶ Masked autoregressive flow (Papamakarios et al., 2017)
- ▶ Glow (Kingma & Dhariwal, 2018)
- ▶ (Block) neural autoregressive flows, spline flows, ... (e.g., Huang et al., 2018; de Cao et al., 2019; Durkan et al., 2019)

Application areas


- ▶ Variational inference in deep generative models
(e.g., Rezende & Mohamed, 2015)
- ▶ Graph neural networks (Liu et al., 2019)
- ▶ Parallel WaveNet (van den Oord et al., 2018)

Application areas

- ▶ Variational inference in deep generative models
(e.g., Rezende & Mohamed, 2015)
- ▶ Graph neural networks (Liu et al., 2019)
- ▶ Parallel WaveNet (van den Oord et al., 2018)
- ▶ Continuous flows
 - ▶ Neural ODEs (e.g, E, 2017; Chen et al., 2018)
 - ▶ Flows on manifolds (e.g., Gemici et al., 2016; Rezende et al., 2020; Mathieu & Nickel, 2020)




From Liu et al. (2019)



From Rezende et al. (2020)

Summary



- ▶ Normalizing flows provide a constructive way to generate rich distributions
- ▶ Key idea: Transform a simple distribution using a flow of successive (invertible) transformations
- ▶ Key ingredient: Change-of-variables trick
- ▶ Jacobians can be computed efficiently, if the transformations are defined appropriately
- ▶ Can be used as a generator and inference mechanism

References

- Chen, T. Q., Rubanova, Y., Bettencourt, J., and Duvenaud, D. K. (2018). Neural Ordinary Differential Equations. In *Advances in Neural Information Processing Systems*.
- De Cao, N., Aziz, W., and Titov, I. (2019). Block Neural Autoregressive Flow. In *Proceedings of the Conference on Uncertainty in Artificial Intelligence*.
- Dinh, L., Krueger, D., and Bengio, Y. (2014). NICE: Non-linear Independent Components Estimation. *arXiv:1410.8516*.
- Dinh, L., Sohl-Dickstein, J., and Bengio, S. (2017). Density Estimation Using Real NVP. In *Proceedings of the International Conference on Learning Representation*.
- Durkan, C., Bekasov, A., Murray, I., and Papamakarios, G. (2019). Neural Spline Flows. In *Advances in Neural Information Processing Systems*.
- E, W. (2017). A Proposal on Machine Learning via Dynamical Systems. *Communications in Mathematics and Statistics*, 5(1):1–11.

References (cont.)

- Gemici, M. C., Rezende, D. J., and Mohamed, S. (2016). Normalizing Flows on Riemannian Manifolds. *arXiv:1611.02304*.
- Huang, C.-W., Krueger, D., Lacoste, A., and Courville, A. (2018). Neural Autoregressive Flows. In *Proceedings of the International Conference on Machine Learning*.
- Kingma, D. P. and Dhariwal, P. (2018). Glow: Generative Flow with Invertible 1×1 Convolutions. In *Advances in Neural Information Processing Systems*.
- Kingma, D. P., Salimans, T., Jozefowicz, R., Chen, X., Sutskever, I., and Welling, M. (2016). Improved Variational Inference with Inverse Autoregressive Flow. In *Advances in Neural Information Processing Systems*.
- Kobyzev, I., Prince, S., and Brubaker, M. (2020). Normalizing Flows: An Introduction and Review of Current Methods. *IEEE Transactions on Pattern Analysis and Machine Intelligence*.
- Liu, J., Kumar, A., Ba, J., Kiros, J., and Swersky, K. (2019). Graph Normalizing Flows. In *Advances in Neural Information Processing Systems*.

References (cont.)

- Mathieu, E. and Nickel, M. (2020). Riemannian Continuous Normalizing Flows. *arXiv:2006.10605*.
- Mohamed, S. and Lakshminarayanan, B. (2016). Learning in Implicit Generative Models. *arXiv:1610.03483*.
- Oord, A., Li, Y., Babuschkin, I., Simonyan, K., Vinyals, O., Kavukcuoglu, K., Driessche, G., Lockhart, E., Cobo, L., Stimberg, F., Casagrande, N., Grewe, D., Noury, S., Dieleman, S., Elsen, E., Kalchbrenner, N., Zen, H., Graves, A., King, H., Walters, T., Belov, D., and Hassabis, D. (2018). Parallel WaveNet: Fast High-Fidelity Speech Synthesis. In *Proceedings of the International Conference on Machine Learning*.
- Papamakarios, G., Nalisnick, E., Rezende, D. J., Mohamed, S., and Lakshminarayanan, B. (2019). Normalizing Flows for Probabilistic Modeling and Inference. *arXiv:1912.02762*.
- Papamakarios, G., Pavlakou, T., and Murray, I. (2017). Masked Autoregressive Flow for Density Estimation. In *Advances in Neural Information Processing Systems*.
- Rezende, D. J. and Mohamed, S. (2015). Variational Inference with Normalizing Flows. In *Proceedings of the International Conference on Machine Learning*.

References (cont.)

- Rezende, D. J., Papamakarios, G., Racanière, S., Albergo, M. S., Kanwar, G., Shanahan, P. E., and Cranmer, K. (2020). Normalizing Flows on Tori and Spheres. In *Proceedings of the International Conference on Machine Learning*.
- Rippel, O. and Adams, R. P. (2013). High-Dimensional Probability Estimation with Deep Density Models. *arXiv:1302.5125*.
- Salvatier, J., Wiecki, T. V., and Fonnesbeck, C. (2016). Probabilistic Programming in Python using PyMC3. *PeerJ Computer Science*, 2:e55.
- Sisson, S. A., Fan, Y., and Tanaka, M. M. (2007). Sequential Monte Carlo Without Likelihoods. *Proceedings of the National Academy of Sciences of the United States of America*, 104(6):1760–1765.
- Tabak, E. G. and Turner, C. V. (2013). A Family of Nonparametric Density Estimation Algorithms. *Communications on Pure and Applied Mathematics*, 66(2):145–164.